Session 13: Stochastic gradient descent

Optimization and Computational Linear Algebra for Data Science

Léo Miolane

Contents

Stochastic gradient descent

Setting

In machine learning, one often has to minimize functions of the form

$$f(x) = \frac{1}{N} \sum_{i=1}^{N} f_i(x).$$

where $f_i : \mathbb{R}^n \to \mathbb{R}$.

Setting

Stochastic gradient descent

$$f(x) = \frac{1}{N} \sum_{i=1}^{N} f_i(x).$$

Starting at some $x_0 \in \mathbb{R}^n$, perform the updates:

$$\begin{aligned} &\text{Pick} \quad i \quad \text{uniformly at random in} \quad \{1, \ldots, N\}, \\ &\text{Update} \quad x_{t+1} = x_t - \alpha_t \nabla f_i(x_t), \end{aligned}$$

Tradeoffs in SGD

Rapidly decaying step sizes

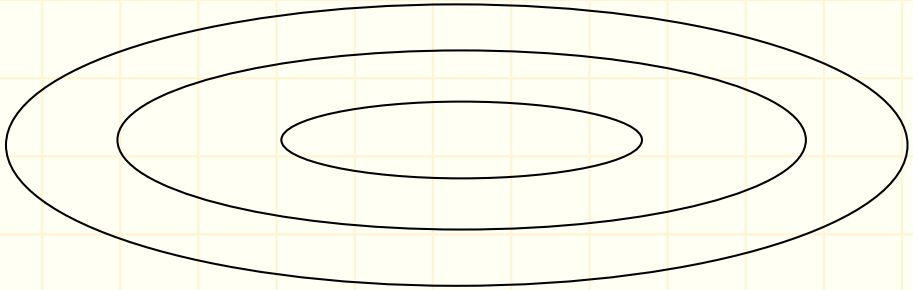

Slowly decaying step sizes

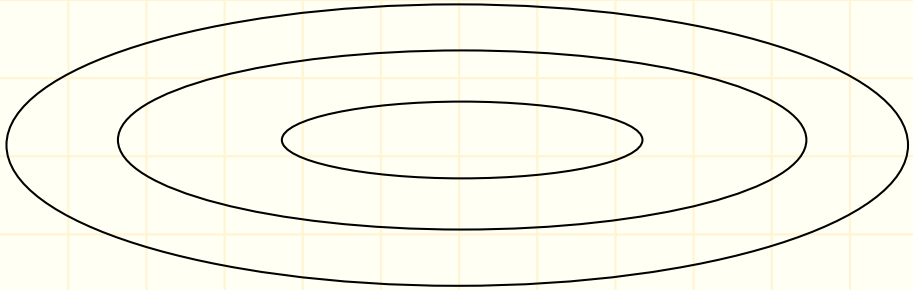

Convergence analysis

Convergence analysis

GD vs SGD

Gradient descent

Stochastic gradient descent

GD vs SGD

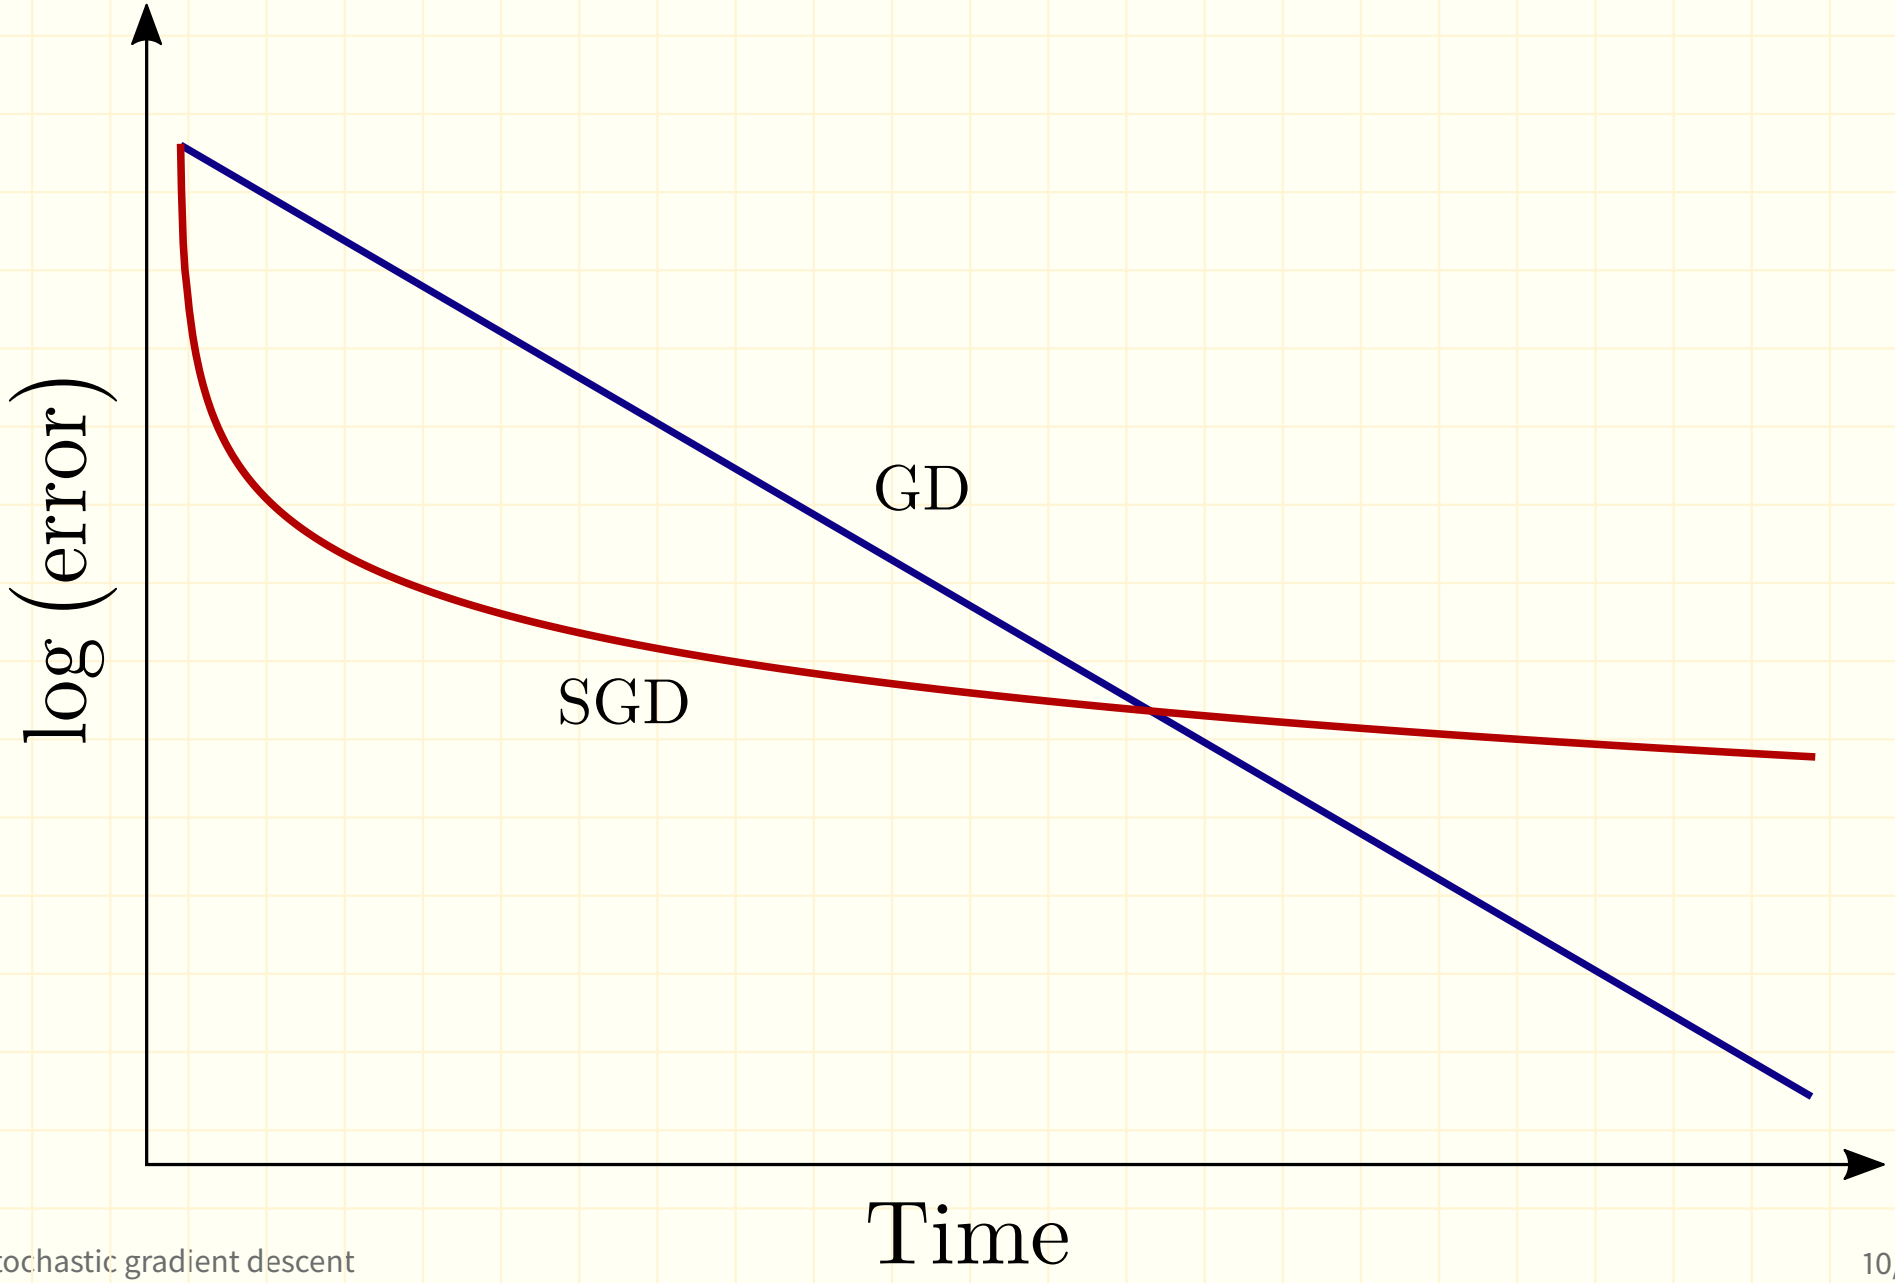

Questions?

Questions?